

Creating a Separate Directory Hash Module

Let's create a separate Python file for the directory hashing functionality that you can import into foo.py.

Here's the implementation:

```
```python
import os
import hashlib

def calculate_directory_hash(directory_path, exclude_dirs=None):
 """
 Calculate a deterministic hash of an entire directory structure.

 Args:
 directory_path (str): Path to the directory to hash
 exclude_dirs (list): List of directory names to exclude (e.g., ['__pycache__', '.git'])

 Returns:
 str: MD5 hash of the directory as a hexadecimal string, or None if directory doesn't exist

 Features:
 - Sensitive to directory structure (folder hierarchy)
 - Sensitive to file contents
 - Insensitive to file access times and other metadata
 """
 if not os.path.exists(directory_path):
 return None

 if exclude_dirs is None:
 exclude_dirs = ['__pycache__', '.git', '.vscode']

 print(f"Calculating hash for directory: {directory_path}")

 # Store both file hashes and structure information
 directory_entries = []

 # Walk through the directory structure
 for root, dirs, files in os.walk(directory_path):
 # Exclude directories we want to skip
 for exclude_dir in exclude_dirs:
 if exclude_dir in dirs:
 dirs.remove(exclude_dir)
```

```

Sort directories for consistent traversal
dirs.sort()

Include directory structure in the hash
rel_dir_path = os.path.relpath(root, directory_path)
if rel_dir_path != '.': # Skip the root directory itself
 # Add entry for this directory to capture the tree structure
 dir_entry = f"DIR:{rel_dir_path}"
 directory_entries.append(dir_entry)
 print(f" - Including directory structure: {rel_dir_path}")

Process each file
for filename in sorted(files): # Sort for consistency
 file_path = os.path.join(root, filename)
 # Skip files we can't read
 if not os.access(file_path, os.R_OK):
 continue

 # Calculate relative path for consistent hashing across systems
 rel_path = os.path.relpath(file_path, directory_path)

 try:
 # Hash the file content
 with open(file_path, 'rb') as f:
 file_content = f.read()

 content_hash = hashlib.md5(file_content).hexdigest()

 # Store file entry with path and hash
 file_entry = f"FILE:{rel_path}:{content_hash}"
 directory_entries.append(file_entry)
 print(f" - Hashed file: {rel_path} -> {content_hash[:8]}...")
 except Exception as e:
 print(f" - Error hashing file {rel_path}: {e}")

Create a combined hash from all entries
if not directory_entries:
 return None

Sort the entries for deterministic order
directory_entries.sort()

Combine all entries into one final hash

```

```
combined_hash = hashlib.md5()
for entry in directory_entries:
 combined_hash.update(entry.encode('utf-8'))
```

```
final_hash = combined_hash.hexdigest()
print(f"Final directory hash: {final_hash}")
return final_hash
```

```
def verify_directory_hash(directory_path, expected_hash, exclude_dirs=None):
 """
```

Verify that a directory's current hash matches the expected hash.

Args:

directory\_path (str): Path to the directory to verify  
expected\_hash (str): Expected hash value  
exclude\_dirs (list): List of directory names to exclude

Returns:

bool: True if the hashes match, False otherwise

```
 """
```

```
 if not expected_hash:
```

```
 print(f"No expected hash provided for {directory_path}, skipping verification")
 return True
```

```
 current_hash = calculate_directory_hash(directory_path, exclude_dirs)
```

```
 if current_hash == expected_hash:
```

```
 print(f"Hash verification PASSED for {directory_path}")
 return True
```

```
 else:
```

```
 print(f"Hash verification FAILED for {directory_path}")
 print(f" Expected: {expected_hash}")
 print(f" Actual: {current_hash}")
 return False
```

```
def generate_hashes_for_directories(directories, exclude_dirs=None):
 """
```

Generate hashes for a list of directories.

Args:

directories (list): List of directory paths to hash  
exclude\_dirs (list): List of directory names to exclude

Returns:

dict: Dictionary mapping directory names to their hash values

```

"""
hashes = {}

for directory in directories:
 if os.path.isdir(directory):
 dir_name = os.path.basename(directory)
 hash_value = calculate_directory_hash(directory, exclude_dirs)
 hashes[dir_name] = hash_value

return hashes
...

```

Now, let's modify your foo.py to import and use this module:

```

```python
import importlib.util
import os
import sys
import directory_hash # Import the directory hash module

def dynamic_import_and_run(script_name, args, verify_hash=False, expected_hash=None):
    # Get the absolute path to the script folder
    script_folder = os.path.join(os.getcwd(), script_name)
    script_path = os.path.join(script_folder, f"{script_name}.py")

    # First check if the script folder exists
    if not os.path.isdir(script_folder):
        result = 1 # Failure
        print(f"Error: Script folder '{script_folder}' not found")
        print(f"=== FAILED {script_name} ({result}) ===\n")
        return result

    # Verify hash if requested
    if verify_hash:
        # Exclude __pycache__ directories by default
        exclude_dirs = ['__pycache__', '.git', '.vscode']
        if not directory_hash.verify_directory_hash(script_folder, expected_hash, exclude_dirs):
            result = 1 # Failure
            print(f"=== FAILED {script_name} ({result}) ===\n")
            return result

    # Check if the actual script file exists
    if not os.path.exists(script_path):
        result = 1 # Failure

```

```

print(f"Error: Script {script_path} not found")
print(f"=== FAILED {script_name} ({result}) ===\n")
return result

```

```

# Load the module
spec = importlib.util.spec_from_file_location(script_name, script_path)
module = importlib.util.module_from_spec(spec)

```

```

try:
    spec.loader.exec_module(module)
except Exception as e:
    result = 1 # Failure
    print(f"Error loading {script_name}: {e}")
    print(f"=== FAILED {script_name} ({result}) ===\n")
    return result

```

```

print(f"\n=== STARTING {script_name} ===")
try:
    func = getattr(module, script_name)
    func_return = func(*args)
    # If we get here, the function completed without errors
    result = 0 # Success
    print(f"Function returned: {func_return}")
    print(f"=== FINISHED {script_name} ({result}) ===\n")
    return result
except Exception as e:
    result = 1 # Failure
    print(f"Error running {script_name}: {e}")
    print(f"=== FAILED {script_name} ({result}) ===\n")
    return result

```

```

def parse_script_string(script_string):
    parts = script_string.split(" and ")
    scripts = []
    for part in parts:
        if ':' in part:
            name, raw_args = part.split(":", 1)
            args = [arg.strip() for arg in raw_args.split(",") if arg.strip()]
        else:
            name = part.strip()
            args = []
        scripts.append((name, args))
    return scripts

```

```

def main():
    if len(sys.argv) < 3:
        print("Usage: python foo.py <log_id> \"A:1,2 and B and C:5\" [--generate-hash]
[--verify-hash]")
        sys.stderr.write("\n")
        sys.exit(1)
        return

    log_id = sys.argv[1]
    script_string = sys.argv[2]

    # Check for hash-related flags
    generate_hash = "--generate-hash" in sys.argv
    verify_hash = "--verify-hash" in sys.argv

    # In-memory hash storage - ADD YOUR KNOWN HASHES HERE
    script_hashes = {
        "A": "1e87a32df7c758c98f3558ea22b7aebd",
        "B": "f91d4068c0e460a54b4bf322fe36805f",
        # Add other script hashes here after generating them
    }

    print(f"Log ID: {log_id}")
    scripts_to_run = parse_script_string(script_string)

    # If in hash generation mode, calculate hashes but don't run scripts
    if generate_hash:
        print("Generating directory hashes...")

        # Create a list of script folders
        script_folders = []
        for name, _ in scripts_to_run:
            script_folder = os.path.join(os.getcwd(), name)
            if os.path.isdir(script_folder):
                script_folders.append(script_folder)

        # Generate hashes for all script folders
        exclude_dirs = ['__pycache__', '.git', '.vscode']
        generated_hashes = directory_hash.generate_hashes_for_directories(script_folders,
exclude_dirs)

        print("Generated hashes:")
        for name, hash_val in generated_hashes.items():
            print(f" {name}: {hash_val}")

```

```

    print("\nCopy these hashes into your script_hashes dictionary for verification:")
    print("script_hashes = {")
    for name, hash_val in generated_hashes.items():
        print(f"    {name}: \"{hash_val}\",")
    print("}")
    return

all_success = True
results = []

for name, args in scripts_to_run:
    # Set up hash verification if needed
    expected_hash = script_hashes.get(name) if verify_hash else None

    # Run the script with optional hash verification
    result = dynamic_import_and_run(name, args, verify_hash, expected_hash)
    results.append({"script": name, "args": args, "result": result})

    # Write each script's result code to stderr
    sys.stderr.write(f"{result}\n")

    if result != 0:
        all_success = False

final_code = 0 if all_success else 1

# Print summary of all results
print("\nScript Results Summary:")
for item in results:
    print(f"    {item['script']}: {item['result']}")

print(f"\nFinal result code to return: {final_code}")

# Write final result code to stderr
sys.stderr.write(f"{final_code}\n")

sys.exit(final_code)

if __name__ == "__main__":
    main()
...

## Benefits of This Modular Approach

```

1. **Reusability**: The `directory_hash.py` module can be imported by any script that needs to verify directory integrity
2. **Maintainability**: Easier to update the hash algorithm or add new features without modifying `foo.py`
3. **Cleaner Code**: Separates concerns - `foo.py` focuses on script execution, `directory_hash.py` handles integrity checking
4. **Testing**: You can test the hash functions independently

Improvements in This Version

1. **Added `exclude_dirs` Parameter**: Default excludes common directories like `__pycache__`, `.git`, and `.vscode`
2. **Added `generate_hashes_for_directories` Function**: A utility to generate hashes for multiple directories at once
3. **Enhanced Documentation**: Better docstrings explaining the purpose and behavior of each function

To use this, save the `directory_hash.py` file in the same folder as `foo.py`, and then you can run `foo.py` as before. The hash generation and verification will now use the external module.

Step-by-Step Explanation of Directory Hash Calculation

Let me walk through the `calculate_directory_hash()` function from the `directory_hash` module, explaining each step with an example to show how it handles directory structures.

Example Directory Structure

Let's use this directory structure for the example:

A/

```
├─ A.py          (content: "def A(): return 'Success'")
├─ data.txt      (content: "Some data")
├─ x/
│   └─ file1.txt (content: "File in X")
│   └─ y/
│       └─ file2.txt (content: "File in Y")
│       └─ z/
│           └─ file3.txt (content: "Deep file")
```

Step 1: Start the Walk Through the Directory

```
for root, dirs, files in os.walk(directory_path):
```

This function walks the directory tree starting at `directory_path`. For each directory, it yields:

- `root`: The current directory path being processed
- `dirs`: List of subdirectories in the current directory
- `files`: List of files in the current directory

For our example, this would process in approximately this order:

1. A (with `dirs=["X"]` and `files=["A.py", "data.txt"]`)
2. A/X/ (with `dirs=["Y"]` and `files=["file1.txt"]`)
3. A/X/Y/ (with `dirs=["Z"]` and `files=["file2.txt"]`)
4. A/X/Y/Z/ (with `dirs=[]` and `files=["file3.txt"]`)

Step 2: Exclude Specified Directories

```
for exclude_dir in exclude_dirs:
```

```
if exclude_dir in dirs:

    dirs.remove(exclude_dir)
```

This removes directories we don't want to include in the hash, like `__pycache__`, `.git`, and `.vscode`. In our example, none of these exist, so nothing is excluded.

Step 3: Sort Directories for Consistent Results

```
dirs.sort()
```

This ensures directories are processed in alphabetical order, regardless of the order returned by the filesystem.

Step 4: Add Directory Structure to Hash

```
rel_dir_path = os.path.relpath(root, directory_path)

if rel_dir_path != '.': # Skip the root directory itself

    # Add entry for this directory to capture the tree structure

    dir_entry = f"DIR:{rel_dir_path}"

    directory_entries.append(dir_entry)

    print(f" - Including directory structure: {rel_dir_path}")
```

This explicitly adds the directory paths (relative to the root) to the entries that will be hashed:

For our example, these directory entries would be added:

- DIR:X
- DIR:X/Y
- DIR:X/Y/Z

By including these entries, the hash becomes sensitive to the directory structure itself, not just its contents.

Step 5: Process Each File

```
for filename in sorted(files):

    file_path = os.path.join(root, filename)

    # Skip files we can't read

    if not os.access(file_path, os.R_OK):
```

```

        continue

    # Calculate relative path for consistent hashing across systems
    rel_path = os.path.relpath(file_path, directory_path)

    try:

        # Hash the file content

        with open(file_path, 'rb') as f:

            file_content = f.read()

            content_hash = hashlib.md5(file_content).hexdigest()

            # Store file entry with path and hash

            file_entry = f"FILE:{rel_path}:{content_hash}"

            directory_entries.append(file_entry)

            print(f" - Hashed file: {rel_path} -> {content_hash[:8]}...")

    except Exception as e:

        print(f" - Error hashing file {rel_path}: {e}")

```

This step:

1. Sorts files alphabetically for consistency
2. Reads each file's content
3. Calculates an MD5 hash of the content
4. Creates an entry with both the file's relative path and content hash
5. Adds this entry to the list of things to hash

For our example, it would add these entries:

- FILE:A.py:md5_of_content_here
- FILE:data.txt:md5_of_content_here
- FILE:X/file1.txt:md5_of_content_here
- FILE:X/Y/file2.txt:md5_of_content_here
- FILE:X/Y/Z/file3.txt:md5_of_content_here

Step 6: Sort All Entries

```
directory_entries.sort()
```

After all directories and files have been processed, all entries are sorted alphabetically. This ensures that the final hash is always the same regardless of the order in which the filesystem returned files and directories.

For our example, after sorting, the entries would be in this order:

- DIR:X
- DIR:X/Y
- DIR:X/Y/Z
- FILE:A.py:hash...
- FILE:data.txt:hash...
- FILE:X/file1.txt:hash...
- FILE:X/Y/file2.txt:hash...
- FILE:X/Y/Z/file3.txt:hash...

Step 7: Generate Final Hash

```
combined_hash = hashlib.md5()

for entry in directory_entries:

    combined_hash.update(entry.encode('utf-8'))

final_hash = combined_hash.hexdigest()
```

Finally, each entry is encoded to bytes and added to the MD5 hash object. The final hexadecimal hash is returned.

How This Captures Directory Structure

Let's see what happens with different structural changes:

Example 1: Moving a File

If we move file2.txt from X/Y/ to X/:

- The old entry FILE:X/Y/file2.txt:hash... is gone
- A new entry FILE:X/file2.txt:hash... is added
- The hash changes due to the path difference

Example 2: Removing a Directory

If we remove the Z directory and move `file3.txt` to Y:

- The directory entry `DIR:X/Y/Z` is gone
- The file entry changes from `FILE:X/Y/Z/file3.txt:hash...` to `FILE:X/Y/file3.txt:hash...`
- The hash changes due to both the structure and path differences

Example 3: Adding a New Directory

If we add a new empty directory X/W:

- A new entry `DIR:X/W` is added
- The hash changes due to the additional directory entry

Summary

The directory hash calculation:

1. Explicitly adds directory paths as separate entries (`DIR:path`)
2. Includes file paths in file entries (`FILE:path:hash`)
3. Sorts everything for consistency
4. Combines all this information into a single hash

This makes the hash sensitive to:

- The exact directory structure
- The location of each file within that structure
- The content of each file

Any change to content, structure, file names, or file locations will result in a different hash.

Understanding Path-Sensitive File Hashing

You're on the right track! Let me clarify exactly how the path of each file affects the hash calculation.

How Files and Paths Are Hashed Together

In the code, **both** the file content **and** its path are included in the hash. Here's the key section that does this:

```
# Calculate relative path for consistent hashing across systems

rel_path = os.path.relpath(file_path, directory_path)

# Hash the file content

with open(file_path, 'rb') as f:
```

```
file_content = f.read()

content_hash = hashlib.md5(file_content).hexdigest()

# Store file entry with path and hash

file_entry = f"FILE:{rel_path}:{content_hash}"

directory_entries.append(file_entry)
```

Let's break this down step by step for your example of moving file2.txt from X/Y/ to X/:

Scenario 1: Before Moving the File

When file2.txt is in X/Y/:

1. The file's relative path is: X/Y/file2.txt
2. The content hash might be: a1b2c3d4... (based on the file contents)
3. The entry added to directory_entries would be:
4. FILE:X/Y/file2.txt:a1b2c3d4...

Scenario 2: After Moving the File

When file2.txt is moved to X/:

1. The file's relative path is now: X/file2.txt
2. The content hash remains: a1b2c3d4... (if the content hasn't changed)
3. The entry added to directory_entries would be:
4. FILE:X/file2.txt:a1b2c3d4...

How This Affects the Final Hash

Later in the code, when the final directory hash is calculated:

```
# Sort the entries for deterministic order

directory_entries.sort()

# Combine all entries into one final hash

combined_hash = hashlib.md5()

for entry in directory_entries:
```

```
combined_hash.update(entry.encode('utf-8'))
```

The entire string `FILE:X/Y/file2.txt:a1b2c3d4...` (before moving) or `FILE:X/file2.txt:a1b2c3d4...` (after moving) is used as input to the hash.

Since these strings are different due to the different paths, they contribute differently to the final hash even though the file content itself hasn't changed.

Why This Matters

This approach ensures that:

1. Moving a file to a different directory changes the hash
2. Renaming a file changes the hash
3. Changing the directory structure changes the hash
4. Changing file content changes the hash

It's actually **not** hashing the file separately from the path - it's combining them into a single entry and then hashing all entries together.

Visual Example

Let's see what the entries might look like before and after moving:

Before moving (file in X/Y/):

DIR:X

DIR:X/Y

FILE:A.py:hash1

FILE:data.txt:hash2

FILE:X/file1.txt:hash3

FILE:X/Y/file2.txt:hash4 ← This entry includes the path X/Y/

After moving (file in X/):

DIR:X

DIR:X/Y

FILE:A.py:hash1

FILE:data.txt:hash2

```
FILE:X/file1.txt:hash3
```

```
FILE:X/file2.txt:hash4 ← This entry has a different path, so it affects  
the hash
```

The file content hash (hash4) remains the same, but the entry string is different because the path changed, which leads to a different final hash.

Step-by-Step Hash Calculation for Directory Structure

I'll walk through exactly how your example directory structure would be processed and hashed using the `directory_hash.py` code. Let's follow the entire process to see how the final hash is formed.

Directory Structure

```
A/
├── A.py          (content: "def A(): return 'Success'")
├── data.txt      (content: "Some data")
├── X/
│   ├── file1.txt (content: "Content of file1")
│   └── Y/
│       └── file2.txt (content: "File in Y")
```

Step 1: Collect Directory Entries

The algorithm walks through the directory tree and collects entries in this order (but we'll sort them later):

Step 1a: Process root directory "A/"

- Files: "A.py" and "data.txt"
- Subdirectories: "X"

For these files, we'd calculate:

```
# Hash "def A(): return 'Success'"
content_hash_A_py = hashlib.md5(b"def A(): return 'Success'").hexdigest()

# Result: "5c5a5db254c3a2ff13edded13db795b0"

# Hash "Some data"
content_hash_data_txt = hashlib.md5(b"Some data").hexdigest()

# Result: "1e50210a0202497fb79bc38b6ade6c34"
```

Add entries:

```
FILE:A.py:5c5a5db254c3a2ff13edded13db795b0
```

```
FILE:data.txt:1e50210a0202497fb79bc38b6ade6c34
```

Step 1b: Process directory "A/X/"

- First, add the directory structure entry:

```
DIR:X
```

- Files: "file1.txt"
- Subdirectories: "Y"

For this file:

```
# Hash "Content of file1"

content_hash_file1_txt = hashlib.md5(b"Content of file1").hexdigest()

# Result: "7d13fd62338f70ddd78fee8497c69304"
```

Add entry:

```
FILE:X/file1.txt:7d13fd62338f70ddd78fee8497c69304
```

Step 1c: Process directory "A/X/Y/"

- First, add the directory structure entry:

```
DIR:X/Y
```

- Files: "file2.txt"
- Subdirectories: None

For this file:

```
# Hash "File in Y"

content_hash_file2_txt = hashlib.md5(b"File in Y").hexdigest()

# Result: "5854c42bed89de16e2fd17fd6c88f5df"
```

Add entry:

```
FILE:X/Y/file2.txt:5854c42bed89de16e2fd17fd6c88f5df
```

Step 2: Sort the Entries

Now we sort all the entries alphabetically:

```
DIR:X
```

```
DIR:X/Y
```

```
FILE:A.py:5c5a5db254c3a2ff13edded13db795b0
```

```
FILE:X/Y/file2.txt:5854c42bed89de16e2fd17fd6c88f5df
```

```
FILE:X/file1.txt:7d13fd62338f70ddd78fee8497c69304
```

```
FILE:data.txt:1e50210a0202497fb79bc38b6ade6c34
```

After sorting (alphabetical order):

```
DIR:X
```

```
DIR:X/Y
```

```
FILE:A.py:5c5a5db254c3a2ff13edded13db795b0
```

```
FILE:data.txt:1e50210a0202497fb79bc38b6ade6c34
```

```
FILE:X/file1.txt:7d13fd62338f70ddd78fee8497c69304
```

```
FILE:X/Y/file2.txt:5854c42bed89de16e2fd17fd6c88f5df
```

Step 3: Calculate the Final Hash

Now, the code combines all these entries into a single MD5 hash:

```
combined_hash = hashlib.md5()
```

```
# For each entry, convert to bytes and update the hash
```

```
combined_hash.update("DIR:X".encode('utf-8'))
```

```
combined_hash.update("DIR:X/Y".encode('utf-8'))
```

```
combined_hash.update("FILE:A.py:5c5a5db254c3a2ff13edded13db795b0".encode('utf-8'))
```

```
combined_hash.update("FILE:data.txt:1e50210a0202497fb79bc38b6ade6c34".encode('utf-8'))
```

```
combined_hash.update("FILE:X/file1.txt:7d13fd62338f70ddd78fee8497c69304".encode('utf-8'))
```

```
combined_hash.update("FILE:X/Y/file2.txt:5854c42bed89de16e2fd17fd6c88f5df"
.encode('utf-8'))

# Get the final hexadecimal hash

final_hash = combined_hash.hexdigest()

# Result would be something like: "a3b7d2e8f1c6509d483e72b9a0f5c64a"
```

How Changes Affect the Hash

Let's see how different changes would alter these entries and thus change the final hash:

1. Moving a file

If we move file2.txt from X/Y/ to X/:

```
- FILE:X/Y/file2.txt:5854c42bed89de16e2fd17fd6c88f5df
+ FILE:X/file2.txt:5854c42bed89de16e2fd17fd6c88f5df
```

This changes one of the entries, producing a different final hash.

2. Removing a directory

If we remove the Y directory completely:

```
- DIR:X/Y
- FILE:X/Y/file2.txt:5854c42bed89de16e2fd17fd6c88f5df
```

Removing these entries changes the input to the hash function, creating a different hash.

3. Modifying file content

If we change the content of file1.txt:

```
- FILE:X/file1.txt:7d13fd62338f70ddd78fee8497c69304
+ FILE:X/file1.txt:2c1743a391305fbf367df8e4f069f9f9
```

The content hash changes, changing the entry, and thus the final hash.

4. Adding a new file

If we add a new file X/Z/newfile.txt:

```
+ DIR:X/Z
```

```
+ FILE:X/Z/newfile.txt:a7c3840dc5a8d46e9a7d4895b77138c1
```

These new entries are included in the hash calculation, changing the final hash.

What Makes This Approach Powerful

1. **Directory Structure is Explicit:** The DIR: entries ensure the folder structure itself is part of the hash
2. **Paths Are Included:** The full relative path of each file is part of its entry
3. **Content is Hashed:** Each file's content hash is included in its entry
4. **Deterministic Ordering:** Sorting entries ensures consistent results regardless of filesystem order
5. **Comprehensive Changes Detection:** Any change to structure or content will produce a different hash

This method creates a fingerprint of the entire directory that captures both its structure and contents, making it perfect for verification.

X-----X

Understanding How Directory Structure Affects Hash Calculation

Great question! You're wondering if strings like "DIR:X/Y" and "DIR:X/Z" might produce the same hash since they contain similar characters. Let me explain why this doesn't happen:

How String-to-Hash Conversion Works

When converting strings to hashes, each character's exact value and position matter. Here's how the encoding and hashing process works:

1. String-to-Bytes Conversion

When we encode a string to UTF-8 bytes:

- Each character is converted to its specific byte representation
- The sequence of bytes preserves the exact order of characters

For example:

- "DIR:X/Y" becomes the bytes: [68, 73, 82, 58, 88, 47, 89]
- "DIR:X/Z" becomes the bytes: [68, 73, 82, 58, 88, 47, 90]

Notice the last byte is different (89 vs 90) because 'Y' and 'Z' are different characters.

2. MD5 Hashing Algorithm

The MD5 hashing algorithm:

- Takes the exact sequence of bytes as input
- Processes them in order
- Generates a hash based on the precise byte values and positions
- Any change, even a single bit, creates a completely different hash

How Our Code Ensures Uniqueness

Let's walk through the specific code that ensures different directory structures produce different hashes:

```
# Sort the entries for deterministic order

directory_entries.sort()

# Combine all entries into one final hash

combined_hash = hashlib.md5()

for entry in directory_entries:

    combined_hash.update(entry.encode('utf-8'))
```

For a structure like A/X/Y vs A/X/Z:

1. We'd have entries like:
 - "DIR:X"
 - "DIR:X/Y" (in the first case)
 - "DIR:X/Z" (in the second case)
2. When we call `entry.encode('utf-8')`, each string is converted to its unique byte representation
3. When we call `combined_hash.update()`, these bytes are processed by the MD5 algorithm
4. Even though "Y" and "Z" are just one character apart, they produce completely different hash outcomes

Example to Demonstrate Uniqueness

Let's see this in action with a simple example:

```
import hashlib

# Two very similar strings

str1 = "DIR:X/Y"

str2 = "DIR:X/Z"

# Hash each string individually
```

```
hash1 = hashlib.md5(str1.encode('utf-8')).hexdigest()

hash2 = hashlib.md5(str2.encode('utf-8')).hexdigest()

print(f"Hash of '{str1}': {hash1}")

print(f"Hash of '{str2}': {hash2}")
```

If you run this code, you'll get completely different hashes:

- Hash of 'DIR:X/Y': a92deba38928ce74b2fbcd35ce1a55b1
- Hash of 'DIR:X/Z': fcf8bc627563c91321abbafa7673b025

This is a fundamental property of cryptographic hash functions - they're designed so that even a tiny change in input produces a completely different output.

Why This Ensures Directory Structure Sensitivity

Because of these properties:

1. Each unique directory structure is encoded as different directory entries
2. Each entry is precisely encoded to bytes
3. The MD5 hash algorithm is sensitive to any small change
4. The final hash will be different for any structural change

So when you move a file from X/Y to X/Z, or rename a directory, the corresponding DIR: entries change, resulting in a completely different final hash, even though the change might seem minor to us.

That's why our hash function correctly detects changes to directory structure - every character in every path matters in the final hash calculation.

A/

```
|— A.py      (content: "def A(): return 'Success'")
|— data.txt  (content: "Some data")
|— X/
|   |— file1.txt (content: "Content of file1")
|   |— Y/
|       |— file2.txt (content: "File in Y")
```